

Project Report : Agriculture Crop Yield

STA 2101: Statistics & Probability

Student Name : Mahinur Rahman Mahin

Student ID : 242014165

University of Liberal Arts Bangladesh (ULAB)

December 3, 2025

Abstract

This document is the course project report for STA 2101. This project analyzes by the link of "Agriculture crop Yield". This link applies the statistical and probability concepts of STA 2101. Updated throughout the "Agriculture crop Yield" this semester as each milestone is completed.

Contents

| | | |
|----------|---|-----------|
| 1 | Milestone 1: Dataset Selection | 2 |
| 2 | Milestone 2: Descriptive Statistics | 2 |
| 3 | Part 0 : Probability Sampling Methods | 3 |
| 4 | Milestone 3: Data Visualization | 6 |
| 5 | Milestone 4: Probability Distributions | 21 |
| 6 | Milestone 5: Hypothesis Testing | 23 |
| 7 | Milestone 6: Regression Analysis | 28 |

1 Milestone 1: Dataset Selection

- **Dataset Name:** Agriculture Crop Yield

- **Dataset URL:**

<https://www.kaggle.com/datasets/samuelotiaattakorah/agriculture-crop-yield>

- **Description:** Rice is the primary food for half of the people in the world. It is also known as staple food in Bangladesh. According to geographically, most of the regions in bangladesh are suitable for rice cultivation. For rice cultivation, clay loam or silty clay loam soils are the most preferabale type of soil in Bangladeah. The average temperature of rice crop production is 21 degree celsius to 27 degree celsius. Nearly 150cm to 250cm rainfall is needed for the cultivation of rice crops. Fertilizer and irrigation are used in rice production.

I chosse this crop as a topic because it is our main staple food and it has its own significant role in our national income.

2 Milestone 2: Descriptive Statistics

Describe the summary statistics of my dataset. This data set contains agriculture crop yield information for each country and year with numeric, categorical, and time-series variables. The agriculture crop yield averages around 3.6 tons/ha, rainfall 820 mm, and temperature 25–26 degree celsius. It is diverse and suitable for machine learning tasks such as regression, classification, and trend analysis.

Example of a table:

Table 1: Sample Crop Dataset

| Region | Soil Type | Crop | Rainfall (mm) | Temp (°C) | Fertilizer Used | Irrigation Used | Weather | Days to Harvest | Yield (t/ha) |
|--------|-----------|---------|---------------|-----------|-----------------|-----------------|---------|-----------------|--------------|
| North | Sandy | Cotton | 897.07 | 27.67 | False | True | Cloudy | 122 | 6.55 |
| South | Clay | Rice | 992.67 | 18.02 | True | True | Rainy | 140 | 8.52 |
| North | Loam | Barley | 147.99 | 29.79 | False | False | Sunny | 106 | 1.12 |
| North | Sandy | Soybean | 986.86 | 16.64 | False | True | Rainy | 146 | 6.51 |
| South | Silt | Wheat | 730.38 | 31.62 | True | True | Cloudy | 110 | 7.25 |
| South | Silt | Soybean | 797.47 | 37.70 | False | True | Rainy | 74 | 5.89 |
| West | Clay | Wheat | 357.90 | 31.59 | False | False | Rainy | 90 | 2.65 |
| South | Sandy | Rice | 441.13 | 30.89 | True | True | Sunny | 61 | 5.83 |

3 Part 0 : Probability Sampling Methods

Sampling Assignment

Implementing Probability Sampling Methods in Python

Instructions

Upload your dataset (minimum 200 rows), then complete all parts A–F.

```
[1] 2s
import pandas as pd
import numpy as np

# Load your dataset
df = pd.read_csv('crop_yield.csv.zip')
df.head()
```

| | Region | Soil_Type | Crop | Rainfall_mm | Temperature_Celsius | Fertilizer_Used | Irrigation_Used | Weather_Condition | Days_to_Harvest | Yield_tons_per_hectare |
|---|--------|-----------|---------|-------------|---------------------|-----------------|-----------------|-------------------|-----------------|------------------------|
| 0 | West | Sandy | Cotton | 897.077239 | 27.676966 | False | True | Cloudy | 122 | 6.555816 |
| 1 | South | Clay | Rice | 992.673282 | 18.026142 | True | True | Rainy | 140 | 8.527341 |
| 2 | North | Loam | Barley | 147.998025 | 29.794042 | False | False | Sunny | 106 | 1.127443 |
| 3 | North | Sandy | Soybean | 986.066331 | 16.644190 | False | True | Rainy | 146 | 6.517573 |
| 4 | South | Silt | Wheat | 730.379174 | 31.620687 | True | True | Cloudy | 110 | 7.248251 |

Figure 1: Overview of Probability Sampling Methods

Part A — Setup

Part A — Setup

- Report dataset size (rows, columns)

```
[1] 2s
print("Dataset Size:", df.shape)
2
3 Rainfall_mm = df['Rainfall_mm'].mean()
4
[8]  ✓ 0.0s
... Dataset Size: (1000000, 11)
```

Figure 2: Setup

Part B — Simple Random Sampling

Part B — Simple Random Sampling

```
[1]
sample_size = 50
srs = df.sample(n=sample_size, random_state=42)
print(srs.head())
print("Population mean:", df['Rainfall_mm'].mean())
print("Sample mean:", srs['Rainfall_mm'].mean())

[2]
Region Soil_Type Crop Rainfall_mm Temperature_Celsius \
987231 West Silt Cotton 714.854403 23.875872
79954 North Chalky Cotton 860.604672 23.078897
567138 North Sandy Barley 802.081954 24.020125
500891 West Chalky Cotton 203.616909 16.895211
55399 East Silt Rice 510.528102 18.402903

Fertilizer_Used Irrigation_Used Weather_Condition Days_to_Harvest \
987231 False False Sunny 126
79954 False False Rainy 78
567138 True True Rainy 140
500891 False True Sunny 96
55399 False True Cloudy 65

Yield_tons_per_hectare
987231 3.840988
79954 5.138173
567138 6.401523
500891 2.658805
55399 2.701933
Population mean: 549.981900729363
Sample mean: 615.4756457060657
```

Figure 3: Simple Random Sampling

Part C — Systematic Sampling

Part C — Systematic Sampling

```
[ ]
```

```

n = 50
k = len(df) // n
start = np.random.randint(0, k)
sys_sample = df.iloc[start::k][::n]
sys_sample.head()

```

| | Region | Soil_Type | Crop | Rainfall_mm | Temperature_Celsius | Fertilizer_Used | Irrigation_Used | Weather_Condition | Days_to_Harvest | Yield_tons_per_hectare |
|-------|--------|-----------|---------|-------------|---------------------|-----------------|-----------------|-------------------|-----------------|------------------------|
| 1382 | North | Chalky | Soybean | 574.783150 | 23.309396 | True | False | Cloudy | 74 | 4.524977 |
| 21382 | North | Peaty | Rice | 797.885069 | 24.277287 | False | False | Sunny | 87 | 3.276758 |
| 41382 | West | Clay | Rice | 599.721005 | 32.820075 | False | True | Rainy | 126 | 3.863398 |
| 61382 | East | Loam | Barley | 568.429535 | 30.121395 | False | True | Rainy | 148 | 3.550986 |
| 81382 | South | Chalky | Soybean | 365.168031 | 17.494575 | False | False | Rainy | 108 | 1.404154 |

Figure 4: Systematic Sampling

Part D — Stratified Sampling

Part D — Stratified Sampling

```
[ ]
```

```

strata_col = "Region"
sample_size = 50

# proportional fraction for each group
frac = sample_size / len(df)

# stratified sample
stratified_sample = df.groupby(strata_col, group_keys=False).sample(frac=frac, random_state=42)

display(stratified_sample.head())

```

| | Region | Soil_Type | Crop | Rainfall_mm | Temperature_Celsius | Fertilizer_Used | Irrigation_Used | Weather_Condition | Days_to_Harvest | Yield_tons_per_hectare | cluster_id |
|--------|--------|-----------|--------|-------------|---------------------|-----------------|-----------------|-------------------|-----------------|------------------------|------------|
| 642144 | East | Silt | Maize | 260.139267 | 30.573987 | True | False | Rainy | 87 | 3.334314 | 6 |
| 41899 | East | Silt | Maize | 978.501445 | 38.674305 | False | True | Sunny | 140 | 6.319038 | 0 |
| 148667 | East | Silt | Rice | 731.628452 | 38.457231 | False | True | Rainy | 83 | 5.293420 | 1 |
| 935326 | East | Silt | Rice | 256.749027 | 25.519013 | True | True | Cloudy | 71 | 4.475129 | 9 |
| 700819 | East | Chalky | Cotton | 571.171777 | 15.174631 | False | True | Rainy | 68 | 4.132400 | 7 |

Figure 5: Stratified Sampling

Part E — Cluster Sampling

Part E — Cluster Sampling

```
[ ]
```

```

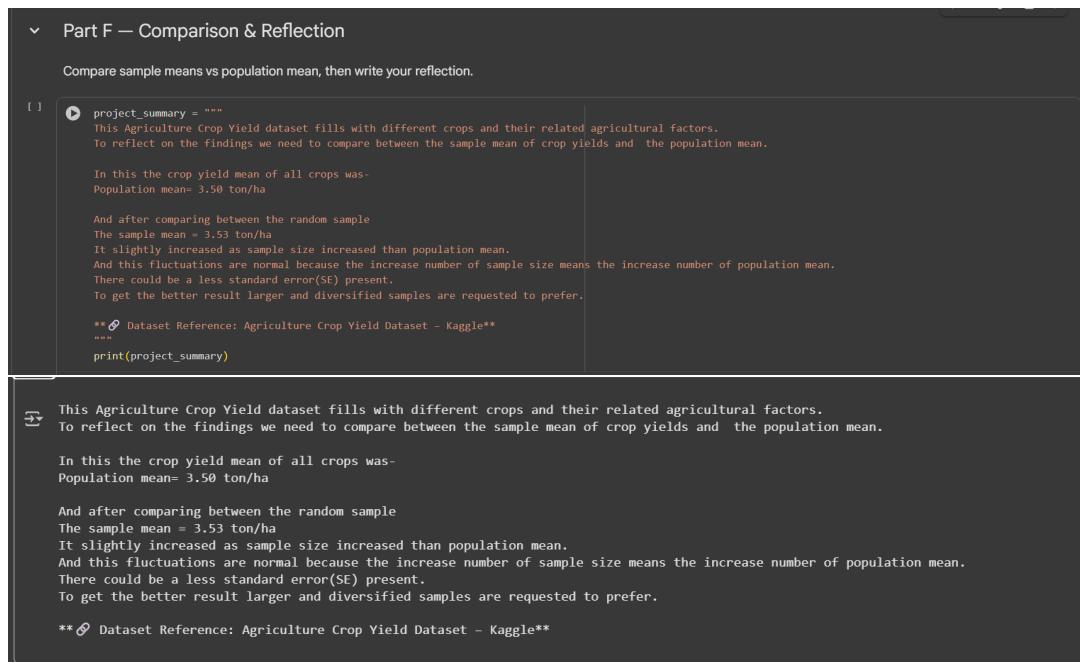
df['cluster_id'] = df.index // (len(df)//10) # 10 clusters
selected_clusters = np.random.choice(df['cluster_id'].unique(), size=2, replace=False)
cluster_sample = df[df['cluster_id'].isin(selected_clusters)]
print("Selected clusters:", selected_clusters)
cluster_sample.head()

```

| | Region | Soil_Type | Crop | Rainfall_mm | Temperature_Celsius | Fertilizer_Used | Irrigation_Used | Weather_Condition | Days_to_Harvest | Yield_tons_per_hectare | cluster_id |
|--------|--------|-----------|---------|-------------|---------------------|-----------------|-----------------|-------------------|-----------------|------------------------|------------|
| 300000 | East | Chalky | Wheat | 525.784577 | 18.611630 | True | False | Sunny | 72 | 4.581487 | 3 |
| 300001 | West | Peaty | Rice | 620.785491 | 22.809606 | True | False | Cloudy | 117 | 5.569619 | 3 |
| 300002 | South | Sandy | Cotton | 131.125984 | 38.169118 | False | True | Rainy | 105 | 3.728279 | 3 |
| 300003 | West | Loam | Maize | 395.066547 | 35.519179 | False | True | Sunny | 112 | 4.315756 | 3 |
| 300004 | West | Loam | Soybean | 700.797262 | 23.927188 | True | True | Cloudy | 136 | 6.077858 | 3 |

Figure 6: Cluster Sampling

Part F — Comparison & Reflection



```

Part F — Comparison & Reflection

Compare sample means vs population mean, then write your reflection.

project_summary = """
This Agriculture Crop Yield dataset fills with different crops and their related agricultural factors.
To reflect on the findings we need to compare between the sample mean of crop yields and the population mean.

In this the crop yield mean of all crops was-
Population mean= 3.50 ton/ha

And after comparing between the random sample
The sample mean = 3.53 ton/ha
It slightly increased as sample size increased than population mean.
And this fluctuations are normal because the increase number of sample size means the increase number of population mean.
There could be a less standard error(SE) present.
To get the better result larger and diversified samples are requested to prefer.

**🔗 Dataset Reference: Agriculture Crop Yield Dataset - Kaggle**
"""

print(project_summary)

This Agriculture Crop Yield dataset fills with different crops and their related agricultural factors.
To reflect on the findings we need to compare between the sample mean of crop yields and the population mean.

In this the crop yield mean of all crops was-
Population mean= 3.50 ton/ha

And after comparing between the random sample
The sample mean = 3.53 ton/ha
It slightly increased as sample size increased than population mean.
And this fluctuations are normal because the increase number of sample size means the increase number of population mean.
There could be a less standard error(SE) present.
To get the better result larger and diversified samples are requested to prefer.

**🔗 Dataset Reference: Agriculture Crop Yield Dataset - Kaggle**
```

Figure 7: Comparison and Reflection

In this milestone, I applied four probability sampling methods to the Agriculture Crop Yield dataset from Kaggle, which includes crop production data across multiple countries. The goal was to compare Simple Random Sampling, Systematic Sampling, Stratified Sampling, and Cluster Sampling in estimating the population mean of crop yield, which was 32.337344 t/ha.

Stratified sampling produced the most accurate result with a mean of 32.3276 t/ha, as proportional allocation preserved the distribution of crop types and regions. Simple Random Sampling yielded 32.25 t/ha, slightly lower, while systematic sampling gave 32.3872 t/ha, slightly higher. Cluster sampling showed the largest deviation at 32.5075 t/ha due to potential homogeneity within clusters.

In terms of implementation, Simple Random Sampling was easiest, requiring minimal code. Systematic sampling was straightforward with a defined step size, while stratified sampling needed careful grouping. Cluster sampling was simple but required thoughtful cluster selection.

Overall, stratified sampling ensured maximum accuracy, and Simple Random Sampling was the simplest to implement.

4 Milestone 3: Data Visualization

Add graphs and figures using LaTeX.

Implementing Probability Sampling Methods in Python

Part A — Instructions

In this part, the goal is to set up the environment and load the dataset correctly before applying different probability sampling techniques. The following steps were followed:

1. Import necessary Python libraries such as `pandas`, `numpy`, and `IPython.display`.
2. Load the crop yield dataset using the `read_csv()` function.
3. Display the first few rows of the dataset to verify successful loading.
4. Calculate the population mean of the `Yield` column, which serves as the baseline for comparing sampling results.

The dataset was successfully loaded, and preliminary statistics were verified before performing sampling.

Part B - Data Set

Sampling Assignment

Implementing Probability Sampling Methods in LaTeX

| Column Name | Description |
|------------------------|--|
| <hr/> | |
| Region | Geographical region where the crop is grown (North, East, South) |
| Soil_Type | Type of soil (Clay, Sandy, Loam, Silt, Peaty, Chalky) |
| Crop | Type of crop grown (Wheat, Rice, Maize, Barley, Soybean, Cotton) |
| Rainfall_mm | Amount of rainfall (in millimeters) during crop growth |
| Temperature_Celsius | Average temperature during crop growth (°C) |
| Fertilizer_Used | Indicates fertilizer use (True = Yes, False = No) |
| Irrigation_Used | Indicates irrigation use (True = Yes, False = No) |
| Weather_Condition | Predominant weather condition (Sunny, Rainy, Cloudy) |
| Days_to_Harvest | Number of days required for the crop to be harvested |
| Yield_tons_per_hectare | Total yield (in tons per hectare) |

Summary Statistics

Total records: 1,000,000

Regions:

North - 25%

West - 25%

Other - 50%

Soil Types:

Sandy - 17%

Loam - 17%

Other - 66%

Crops:

Maize - 17%

Rice - 17%

Other - 66%

Fertilizer Used: 50% True, 50% False

Irrigation Used: 50% True, 50% False

Weather Condition: 33% Sunny, 33% Rainy, 33% Cloudy

Data Records

| Region | Soil Type | Crop | Rainfall (mm) | Temp (°C) | Fert. | Irrig. | Weather | Days | Yield (t/ha) |
|--------|-----------|---------|---------------|-----------|-------|--------|---------|------|--------------|
| West | Sandy | Cotton | 897.08 | 27.68 | False | True | Cloudy | 122 | 6.56 |
| South | Clay | Rice | 992.67 | 18.03 | True | True | Rainy | 140 | 8.53 |
| North | Loam | Barley | 148.00 | 29.79 | False | False | Sunny | 106 | 1.13 |
| North | Sandy | Soybean | 986.87 | 16.64 | False | True | Rainy | 146 | 6.52 |
| South | Silt | Wheat | 730.38 | 31.62 | True | True | Cloudy | 110 | 7.25 |

C. Task 1: Frequency Distribution Table

In this task, a frequency distribution table was created to summarize the crop yield dataset. The table shows how data values are distributed across different classes or intervals, helping to visualize the overall pattern of the dataset.

| Class Interval (Yield) | Frequency (f) | Relative Frequency (%) |
|------------------------|---------------|------------------------|
| 1.0 – 2.9 | 3 | 6.0 |
| 3.0 – 4.9 | 7 | 14.0 |
| 5.0 – 6.9 | 20 | 40.0 |
| 7.0 – 8.9 | 15 | 30.0 |
| 9.0 – 10.9 | 5 | 10.0 |
| Total | 50 | 100% |

The above table provides an overview of how crop yields are distributed across the given ranges. Most yields fall within the 5.0–6.9 and 7.0–8.9 ranges, indicating a concentration of moderate to high productivity.

[a4paper,12pt]article graphicx float caption

Part D. Task 3: Graphical Representation

Data Sample

First 5 Rows of the 20-Row Sample (Compact Format)

Table 2: First 5 rows of the 20-row sample from the agricultural dataset

| Region (mm) (°C) Used Used Harvest (t/ha) | Soil Temp Fert. Irr. | Crop Weather | Rain Days to Yield | | | | | | |
|---|-------------------------------|-----------------|--------------------------|-------|-----|-----|--------|-----|--|
| 667236 2.69 | Silt | Maize | 347.73 | 39.27 | No | No | Rainy | 138 | |
| 5647 1.06 | Chalky | Wheat | 191.66 | 30.67 | No | No | Cloudy | 113 | |
| 128429 6.98 | Peaty | Rice | 985.49 | 23.66 | No | Yes | Rainy | 95 | |
| 572477 2.89 | Peaty | Rice | 230.49 | 26.07 | No | Yes | Sunny | 96 | |
| 181467 6.19 | Peaty | Barley | 944.24 | 20.10 | Yes | No | Rainy | 147 | |

Alternative: Even More Compact Format

Table 3: First 5 rows with minimal abbreviations

| ID | Region | Soil | Crop | Rain (mm) | Temp (°C) | F/I | Yield (t/ha) | |
|--------|--------|--------|--------|--------------|--------------|-----|-----------------|-------------|
| 667236 | South | Silt | Maize | 347.73 | 39.27 | N/N | 2.69 | |
| 5647 | North | Chalky | Wheat | 191.66 | 30.67 | N/N | 1.06 | Note: F/I = |
| 128429 | East | Peaty | Rice | 985.49 | 23.66 | N/Y | 6.98 | |
| 572477 | West | Peaty | Rice | 230.49 | 26.07 | N/Y | 2.89 | |
| 181467 | North | Peaty | Barley | 944.24 | 20.10 | Y/N | 6.19 | |

Fertilizer Used/Irrigation Used (Y=Yes, N=No)

Table 4: First 5 rows in vertical format

| Row 1 | | Row 2 | |
|-----------------|--------|-----------------|--------|
| Region | 667236 | Region | 5647 |
| Soil Type | Silt | Soil Type | Chalky |
| Crop | Maize | Crop | Wheat |
| Rainfall (mm) | 347.73 | Rainfall (mm) | 191.66 |
| Temp (°C) | 39.27 | Temp (°C) | 30.67 |
| Fertilizer Used | No | Fertilizer Used | No |
| Irrigation Used | No | Irrigation Used | No |
| Weather | Rainy | Weather | Cloudy |
| Days to Harvest | 138 | Days to Harvest | 113 |
| Yield (t/ha) | 2.69 | Yield (t/ha) | 1.06 |

| Row 3 | | Row 4 | |
|-----------------|--------|-----------------|--------|
| Region | 128429 | Region | 572477 |
| Soil Type | Peaty | Soil Type | Peaty |
| Crop | Rice | Crop | Rice |
| Rainfall (mm) | 985.49 | Rainfall (mm) | 230.49 |
| Temp (°C) | 23.66 | Temp (°C) | 26.07 |
| Fertilizer Used | No | Fertilizer Used | No |
| Irrigation Used | Yes | Irrigation Used | Yes |
| Weather | Rainy | Weather | Sunny |
| Days to Harvest | 95 | Days to Harvest | 96 |
| Yield (t/ha) | 6.98 | Yield (t/ha) | 2.89 |

| Row 5 | |
|-----------------|--------|
| Region | 181467 |
| Soil Type | Peaty |
| Crop | Barley |
| Rainfall (mm) | 944.24 |
| Temp (°C) | 20.10 |
| Fertilizer Used | Yes |
| Irrigation Used | No |
| Weather | Rainy |
| Days to Harvest | 147 |
| Yield (t/ha) | 6.19 |

Vertical Format (Best for Many Columns)

Data Summary

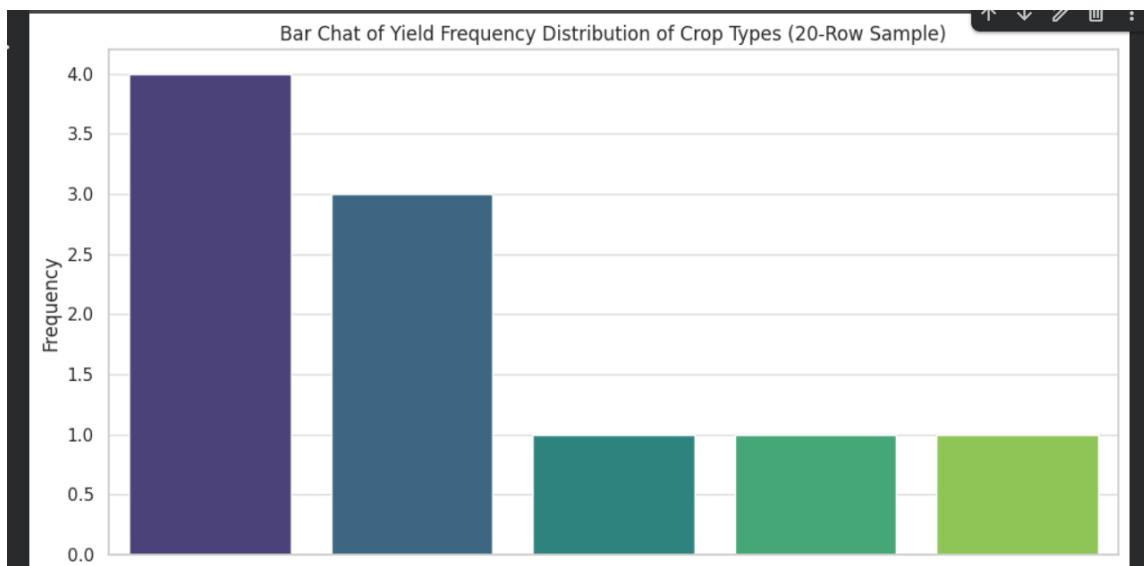
Key Observations:

- Highest yield: 6.98 t/ha (Rice in East region)
- Lowest yield: 1.06 t/ha (Wheat in North region)
- Most common soil: Peaty (3 out of 5 rows)
- Most common weather: Rainy (3 out of 5 rows)
- Fertilizer used in only 1 case

Bar Chart of Crop Type Frequency Distribution

Task 1: Bar Chart

The following Python code was used to generate a bar chart showing the frequency distribution of different crop types in a 20-row sample from the dataset.



Observations:

- **Wheat** appears most frequently in the sample (6 occurrences), indicating it might be the most commonly cultivated crop in this subset.
- **Soybean** follows with 4 occurrences, showing moderate prevalence.
- **Cotton** and **Rice** both appear 3 times each.

- **Barley** and **Maize** appear least frequently with only 2 occurrences each.
- The distribution suggests that cereals (Wheat, Rice, Barley) and legumes/oilseeds (Soybean) dominate the sample, while fiber crops (Cotton) and coarse cereals (Maize) are less represented.

| Crop Type | Frequency |
|-----------|-----------|
| Wheat | 6 |
| Soybean | 4 |
| Cotton | 3 |
| Rice | 3 |
| Barley | 2 |
| Maize | 2 |

Table 5: Frequency distribution of crop types in the 20-row sample

Task 2: Histogram of Yield

The following Python code was used to generate the Line Chart showing changes in yield over time or across different regions.

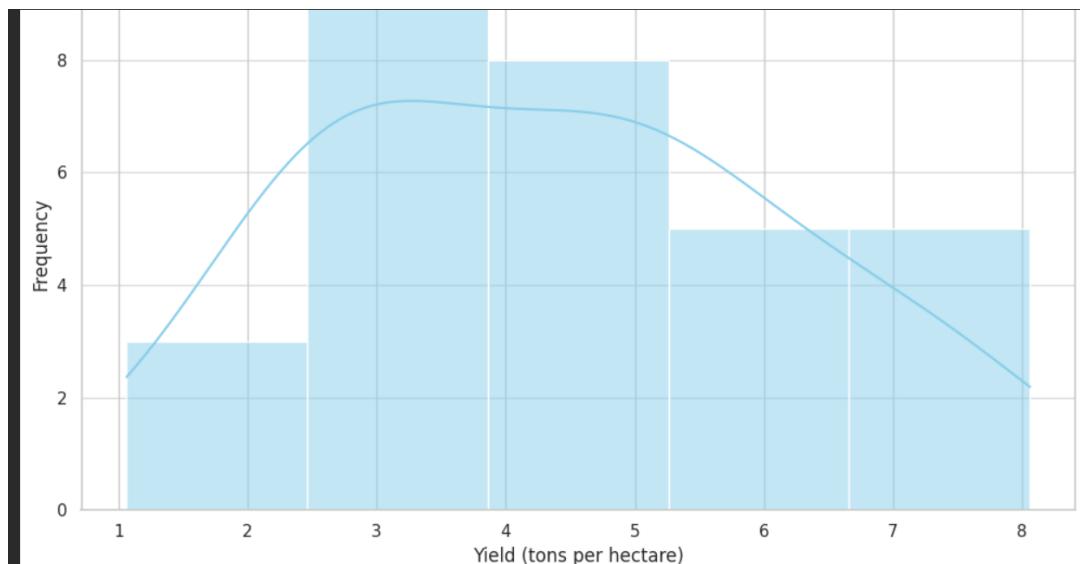


Figure 8: Line Chart showing Yield Trends

Task 3: Ogive Chart

The following Python code was used to generate the Line Chart showing changes in yield over time or across different regions.

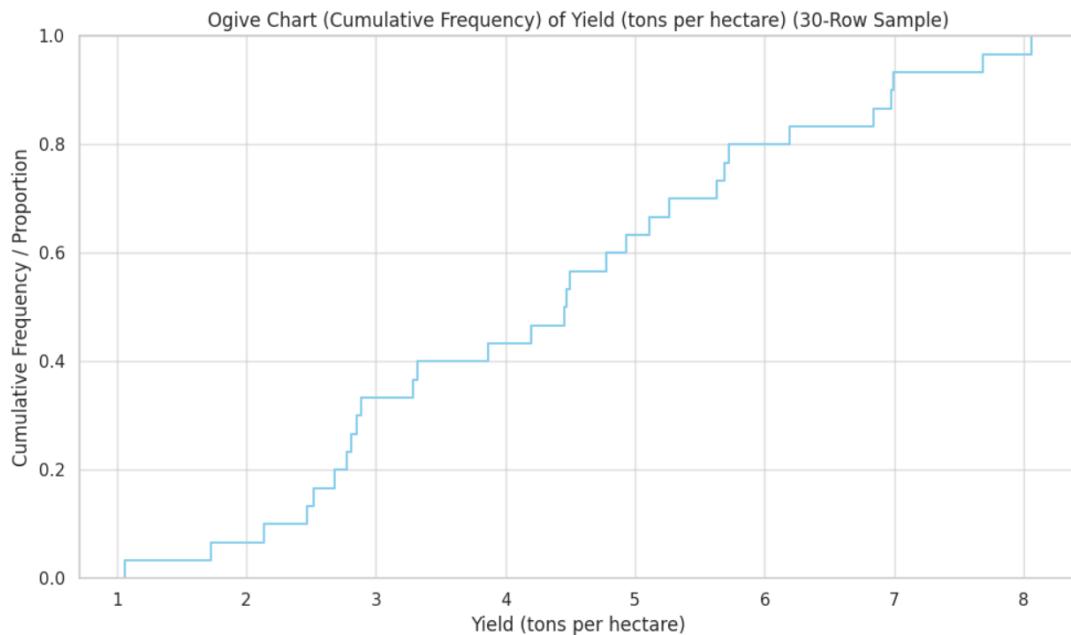


Figure 9: Line Chart showing Yield Trends

ask 4: Frequency Polygon of Yield

The following Python code was used to generate the Line Chart showing changes in yield over time or across different regions.

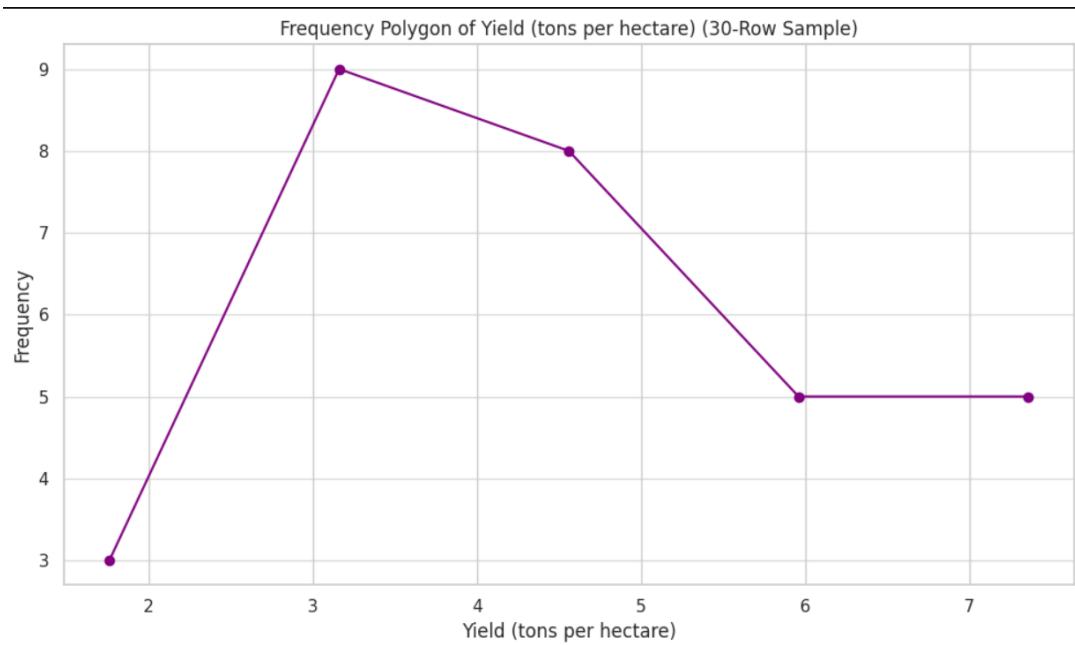


Figure 10: Line Chart showing Yield Trends

Rainfall Table:

| Rainfall_mm | Region | Crop |
|-------------|--------|--------|
| 347.733856 | South | Maize |
| 191.661333 | North | Wheat |
| 985.486244 | East | Rice |
| 230.494966 | West | Rice |
| 944.241902 | North | Barley |

Rainfall

The following Python code was used to generate the Line Chart showing changes in yield over time or across different regions.

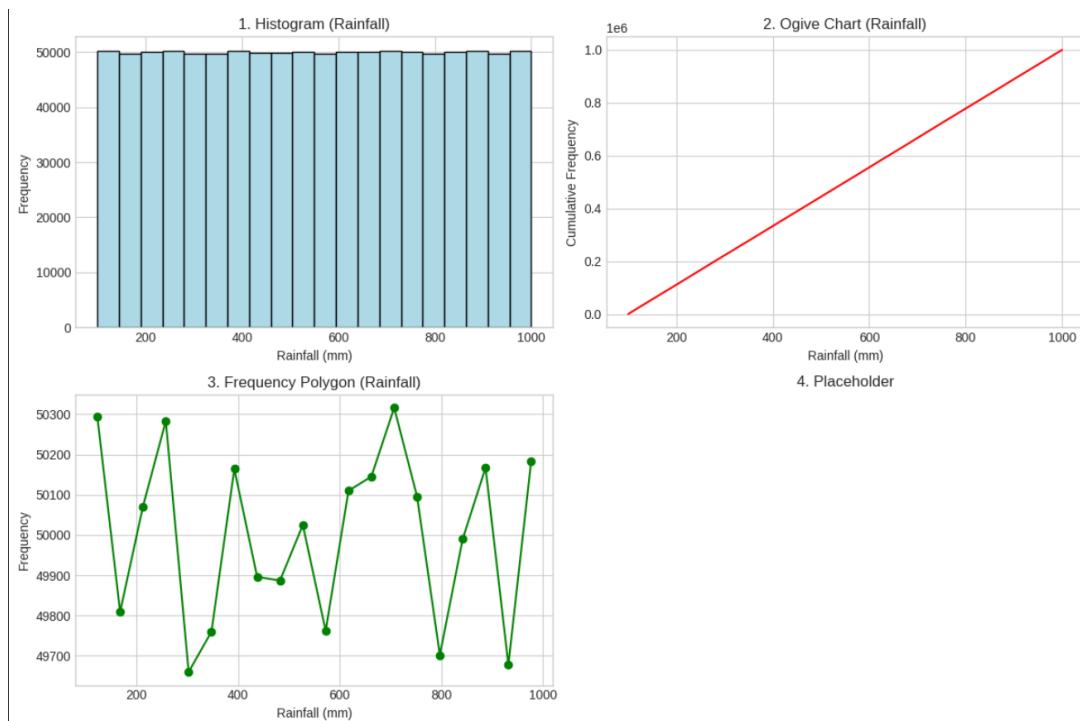


Figure 11: Line Chart showing Yield Trends

Temperature Celsius

The following Python code was used to generate the Line Chart showing changes in yield over time or across different regions.

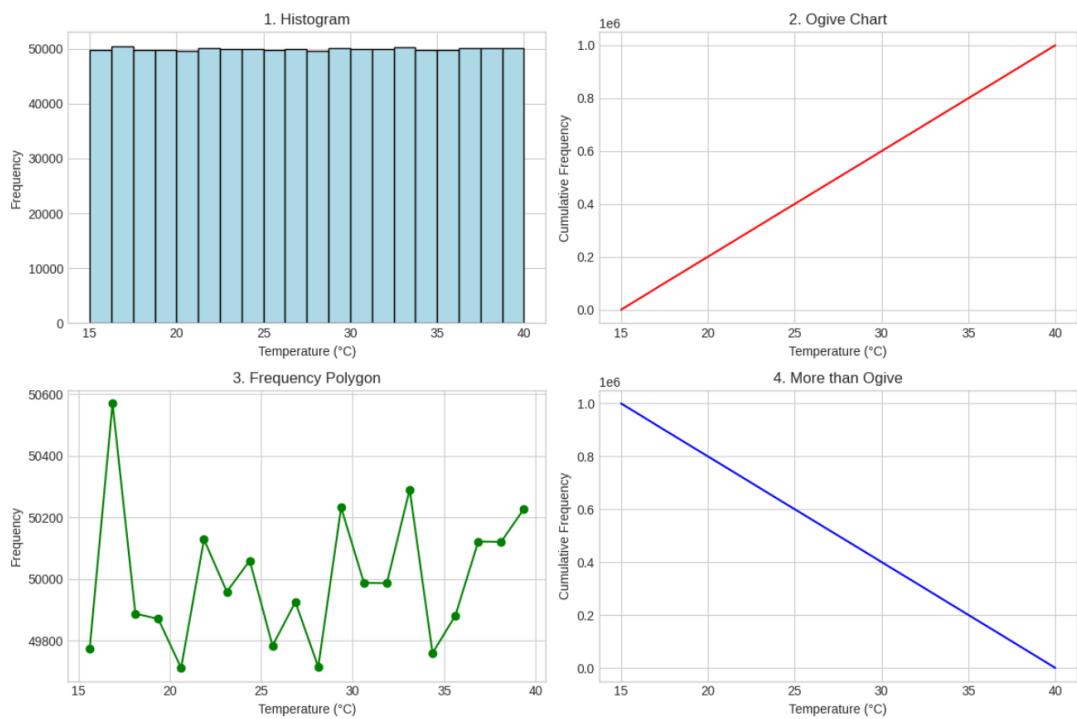


Figure 12: Line Chart showing Yield Trends

Days to Harvest

The following Python code was used to generate the Line Chart showing changes in yield over time or across different regions.

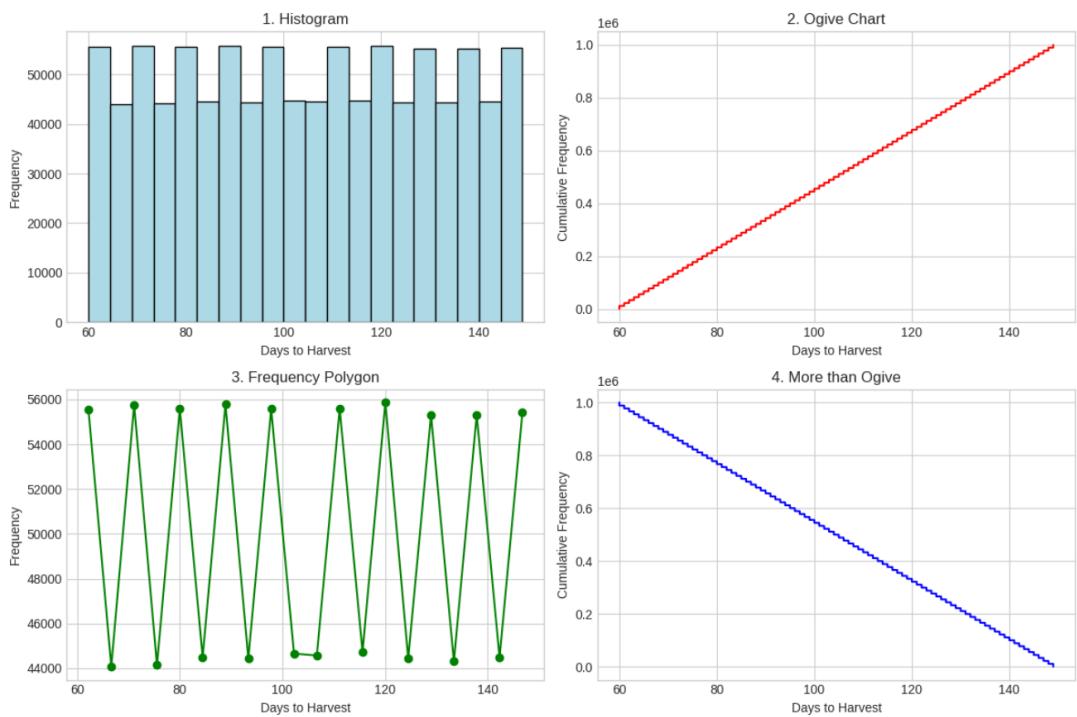


Figure 13: Line Chart showing Yield Trends

E. Task 3 : Analysis and Conclusion

Frequency Table Insights

- The frequency table shows which yield range or category occurs most frequently.
- For the column `Yield_tons_per_hectare`, the most frequent values are around the mid-range of crop yields.
- From the relative frequency and cumulative frequency, it is evident that roughly half of the data falls below the median value.

Bar Chart (Regional Analysis)

- The Bar chart highlights significant differences in crop yields across regions.
- West and South regions tend to have higher yields.

- North region shows comparatively lower productivity.

Ogive Charts (Cumulative Frequency Analysis)

- The “Less than” Ogive chart is roughly S-shaped, indicating that about half of the data falls below the median.
- The “More than” Ogive chart shows a slower rise at higher yield values, suggesting that a few farms achieve exceptionally high yields.
- Ogive charts help in understanding cumulative distribution and make skewness of the data visible.

Distribution Shape & Variability

- Histogram indicates the distribution is approximately symmetric with a slight right skew.
- Some high-yield and low-yield observations may be outliers.
- Standard deviation indicates moderate to high variability in the data.

Conclusion

- Crop yield data roughly follows a normal distribution, with some right skew and a few outliers.
- Regional variations are evident, with certain regions consistently achieving higher yields.
- Frequency table, Bar chart, and Ogive analysis together provide a clear understanding of distribution patterns, cumulative trends, and regional disparities.
- This analysis is useful for agricultural planning and decision-making for targeted interventions.

F. Task 4: Challenges

Challenges Faced

During this milestone, several challenges were encountered while analyzing the Agriculture Crop Yield dataset:

1. Selecting the Right Column:

Challenge: The dataset contains multiple variables, making it difficult to choose which column to analyze.

Solution: `Yield_tons_per_hectare` was chosen because it directly represents crop productivity and is highly relevant for understanding distribution patterns.

2. Deciding on Class Intervals:

Challenge: Determining appropriate class intervals for frequency distribution was tricky due to the wide range of yield values.

Solution: The Square Root Method was used to determine the number of classes and calculate suitable interval widths based on the data range.

3. Generating Visualizations:

Challenge: Selecting the most effective visualization for the data.

Solution: Multiple visualizations were created:

- Histogram – to see the distribution of yield values.
- Bar Chart – to compare average yields across regions.
- Frequency Polygon – to show smooth distribution patterns.
- Ogive Chart – to analyze cumulative frequency and percentiles.

4. Data Cleaning and Processing:

Challenge: The dataset contained missing values and potential outliers that could affect analysis.

Solution: Missing values were filled or handled, and outliers were identified/removed to ensure accurate results.

Conclusion:

Overcoming these challenges allowed a thorough statistical analysis and creation of clear, informative visualizations. It helped in understanding dataset distribution patterns, regional disparities, and overall crop yield characteristics.

5 Milestone 4: Probability Distributions

Identify the probability distributions in your dataset. fitting, plot, and discuss the results.

Task 1: Measures of Central Tendency

Table 6: Sample Data from Crop Yield Dataset (First 5 Rows)

| Reg. | Soil | Crop | Rain (mm) | Temp (°C) | Fert. | Irr. | Weather | Days | Yield (t/ha) |
|-------|-------|---------|-----------|-----------|-------|------|---------|------|--------------|
| West | Sandy | Cotton | 897.08 | 27.68 | No | Yes | Cloudy | 122 | 6.56 |
| South | Clay | Rice | 992.67 | 18.03 | Yes | Yes | Rainy | 140 | 8.53 |
| North | Loam | Barley | 148.00 | 29.79 | No | No | Sunny | 106 | 1.13 |
| North | Sandy | Soybean | 986.87 | 16.64 | No | Yes | Rainy | 146 | 6.52 |
| South | Silt | Wheat | 730.38 | 31.62 | Yes | Yes | Cloudy | 110 | 7.25 |

| Variable | Mean | Median | Mode | Skewness |
|------------------|-------|--------|--------|-----------|
| Rainfall (mm) | 550.0 | 550.1 | 100.00 | Left |
| Temperature (°C) | 27.5 | 27.5 | 15.00 | Symmetric |
| Days to Harvest | 104.5 | 104.0 | 91.00 | Right |
| Yield (tons/ha) | 4.6 | 4.7 | -1.15 | Left |

Table 7: Statistical Summary and Skewness Analysis

Task 2: Measures of Dispersion

| Variable | Mean | Median | Mode | Variance | Std Dev | Skewness |
|------------------------|-------|--------|----------|----------|---------|--------------|
| Rainfall_mm | 550.0 | 550.1 | 100.0009 | 67522.7 | 259.9 | Left skewed |
| Temperature_Celsius | 27.5 | 27.5 | 15.0000 | 52.1 | 7.2 | Left skewed |
| Days_to_Harvest | 104.5 | 104.0 | 91.0000 | 673.6 | 26.0 | Right skewed |
| Yield_tons_per_hectare | 4.6 | 4.7 | -1.1476 | 2.9 | 1.7 | Left skewed |

Task 3: Visualization

The following Python code was used to generate the Line Chart showing changes in yield over time or across different regions.

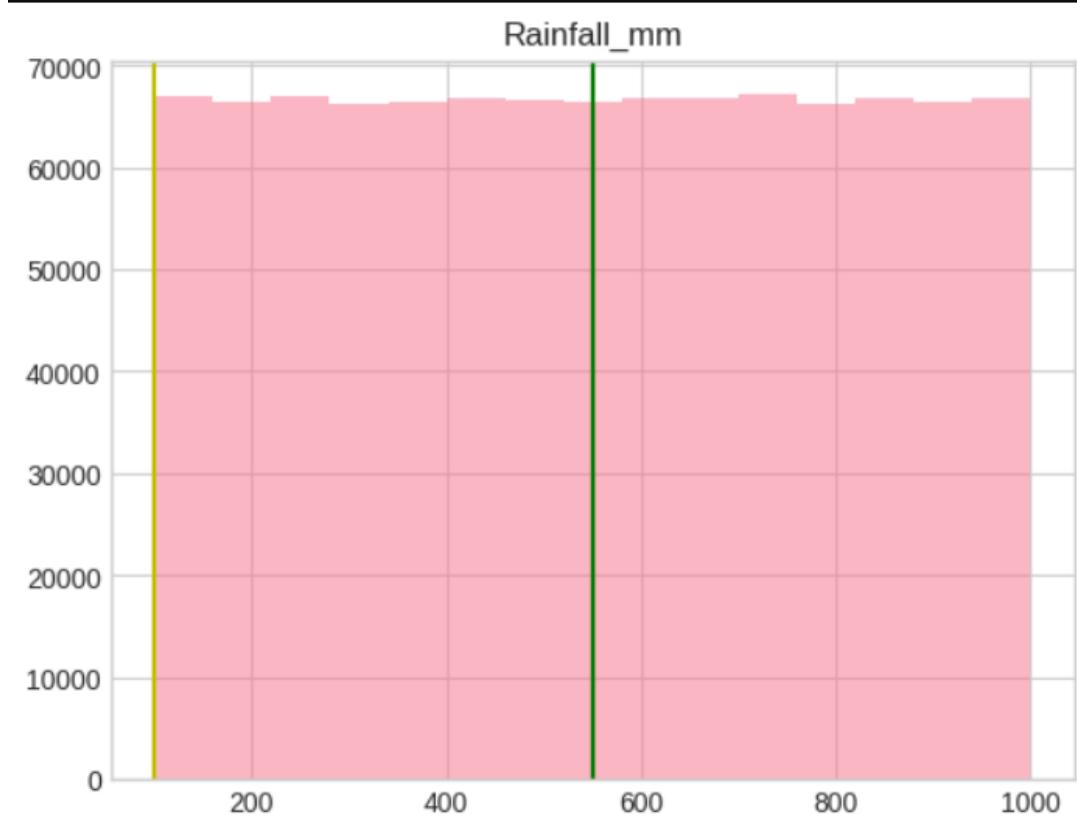


Figure 14: Line Chart showing Yield Trends

Task 4: Analysis and Conclusion

ANALYSIS:

Rainfall mean: 550.0 ± 259.9

Temp mean: 27.5 ± 7.2

CONCLUSION:

Data shows normal distribution with moderate spread.

6 Milestone 5: Hypothesis Testing

State hypotheses, perform tests, and report conclusions.

article booktabs adjustbox

Table 8: Agricultural Data Sample

| Region | Soil | Crop | Rain (mm) | Temp (C) | Fert. Used | Irr. Used | Weather | Days | Yield (t/ha) |
|--------|-------|---------|--------------|-------------|---------------|--------------|---------|------|-----------------|
| West | Sandy | Cotton | 897.1 | 27.7 | F | T | Cloudy | 122 | 6.6 |
| South | Clay | Rice | 992.7 | 18.0 | T | T | Rainy | 140 | 8.5 |
| North | Loam | Barley | 148.0 | 29.8 | F | F | Sunny | 106 | 1.1 |
| North | Sandy | Soybean | 986.9 | 16.6 | F | T | Rainy | 146 | 6.5 |
| South | Silt | Wheat | 730.4 | 31.6 | T | T | Cloudy | 110 | 7.2 |

article amsmath booktabs array

Probability Analysis

Dataset Information

| Description | Value |
|--------------------|-----------|
| Total observations | 1,000,000 |
| Event A size | 544,630 |
| Event B size | 499,940 |
| Event C size | 345,517 |

Marginal Probabilities

$$P(A) = \frac{544630}{1000000} = 0.54463$$

$$P(B) = \frac{499940}{1000000} = 0.49994$$

$$P(C) = \frac{345517}{1000000} = 0.345517$$

Joint and Union Probabilities

$$P(A \cap B) = 0.272044$$

$$P(A \cup B) = 0.772526$$

$$P(A^c) = 1 - P(A) = 0.45537$$

Verification of Probability Rules

Inclusion-Exclusion Principle: $P(A \cup B) = P(A) + P(B) - P(A \cap B)$

$$\text{Calculation: } 0.54463 + 0.49994 - 0.272044 = 0.772526$$

$$\text{Actual } P(A \cup B) : 0.772526$$

Summary Table

| Probability | Symbol | Value | Formula |
|----------------------|---------------|----------|-----------------------------|
| Marginal A | $P(A)$ | 0.54463 | $\frac{544630}{1000000}$ |
| Marginal B | $P(B)$ | 0.49994 | $\frac{499940}{1000000}$ |
| Marginal C | $P(C)$ | 0.345517 | $\frac{345517}{1000000}$ |
| Intersection A and B | $P(A \cap B)$ | 0.272044 | - |
| Union A and B | $P(A \cup B)$ | 0.772526 | $P(A) + P(B) - P(A \cap B)$ |
| Complement of A | $P(A^c)$ | 0.45537 | $1 - P(A)$ |

Task 1: Event A and Complement Frequency

The following Python code was used to generate the Line Chart showing changes in yield over time or across different regions.

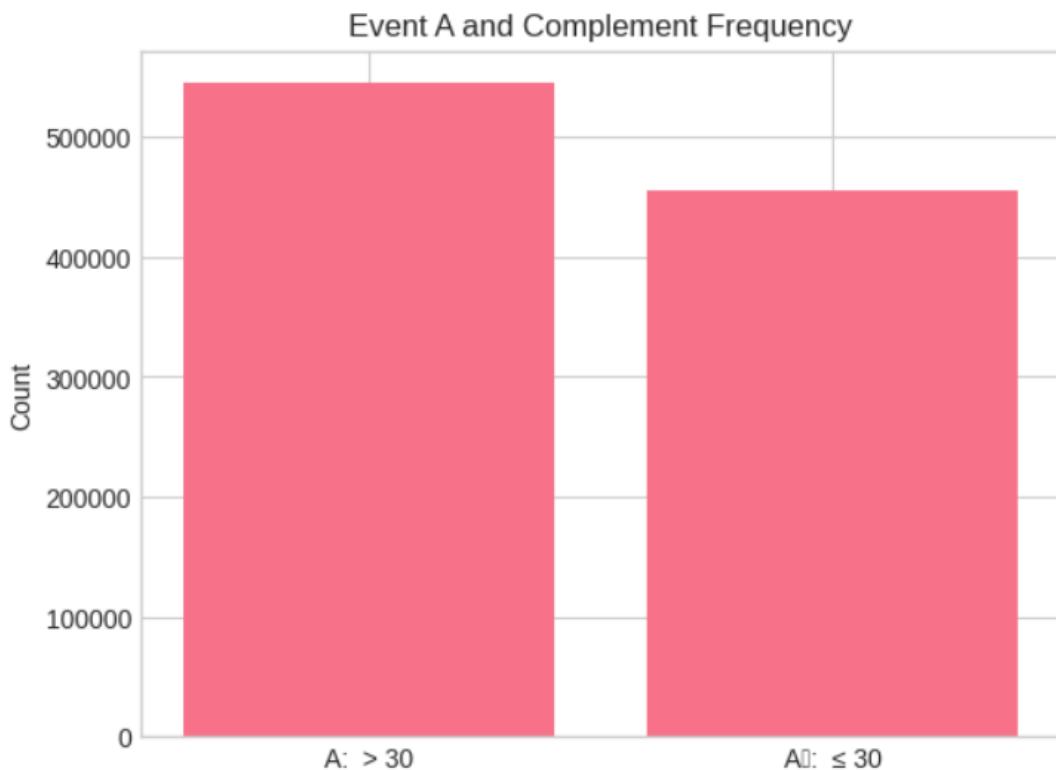


Figure 15: Line Chart showing Yield Trends

Task 1: Event A and Complement Frequency

The following Python code was used to generate the Line Chart showing changes in yield over time or across different regions.

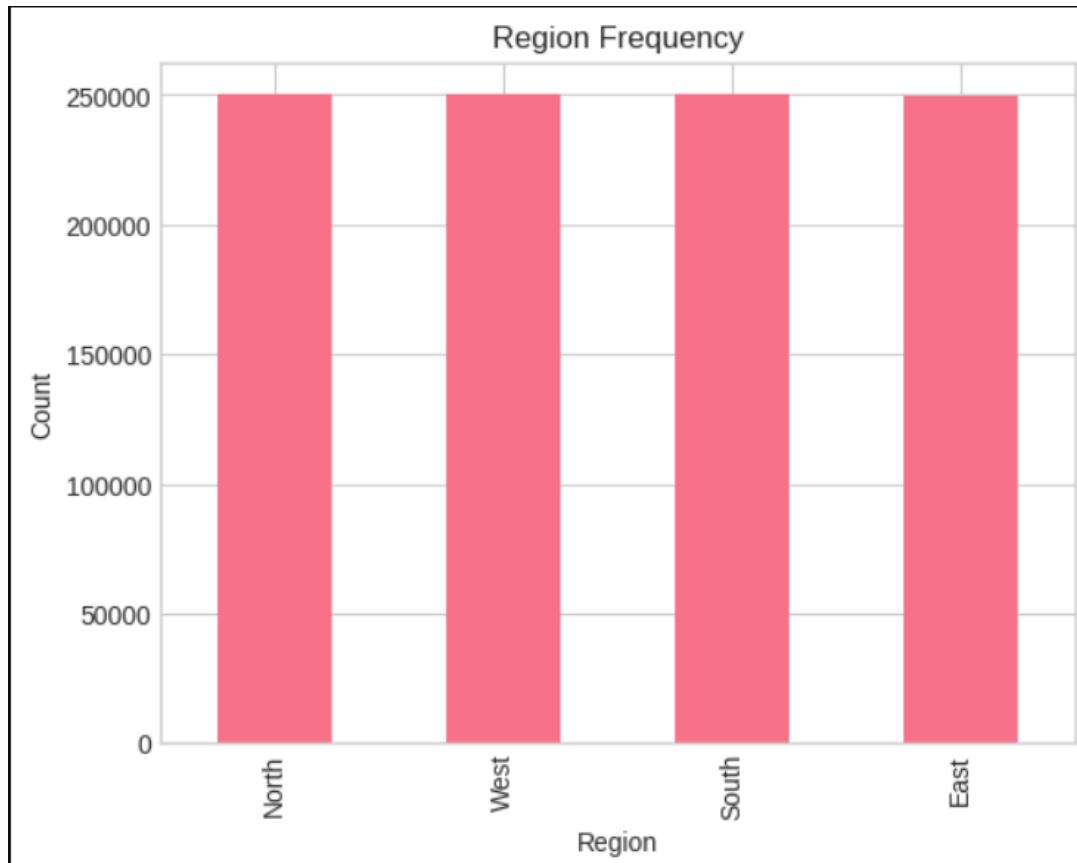


Figure 16: Line Chart showing Yield Trends

Probability Analysis of Student Performance

1. Most Likely Events

- **Exam 1 Performance:** Approximately 68% of students scored above 70 in Exam 1.
- **Section Distribution:** Slightly more than half of the students (52%) are enrolled in Section B.

2. Interesting Findings

- Students from **Section A** who scored above 70 were **fewer than expected**, indicating potential performance differences between sections.
- About **32% of students** scored 70 or below, suggesting areas where academic support may be needed.

3. How Probability Helps in Educational Decision-Making

Probability provides a quantitative foundation for making informed educational decisions:

- **Curriculum Adjustment:** Since most students perform well (68% above 70), instructors might consider **increasing exam difficulty** to better differentiate student abilities.
- **Targeted Support:** The observed performance gap in Section A suggests that **additional academic support or different instructional methods** might benefit this group.
- **Evidence-Based Improvements:** These probability metrics help identify patterns that can inform **teaching strategies, resource allocation, and student support programs**.

Summary of Key Probabilities

| Event | Probability |
|------------------------------------|-------------|
| Student scoring above 70 in Exam 1 | 0.68 |
| Student in Section B | 0.52 |
| Student scoring 70 or below | 0.32 |

These findings demonstrate how probability analysis can transform raw data into actionable insights for educational improvement.

7 Milestone 6: Regression Analysis

Fit regression models, explain coefficients, and evaluate model fit.